

Atlassian – Senior Data Engineer Interview Master Cheat Sheet

Python String Manipulation, JSON Parsing, Scalability Reasoning, and One Unified Interview Code Example

1. Senior Data Engineer – What Interviewers Look For

- Clean, defensive Python code
- Handling messy real-world input
- Strong JSON manipulation skills
- Performance awareness ($O(n)$, memory)
- Ability to explain scaling to Spark/Databricks

2. One End-to-End Python Example (Covers 90% Interview Needs)

Scenario:

You receive raw string input representing event logs (API / Kafka style).

Tasks:

- Clean strings
- Parse JSON
- Handle missing/malformed fields
- Aggregate metrics
- Produce final JSON output

```
import json
from collections import defaultdict

def process_events(raw_lines):
    results = defaultdict(int)

    for line in raw_lines:
        # Step 1: Clean string input
        line = line.strip()
        if not line:
            continue

        # Step 2: Safe JSON parsing
        try:
            event = json.loads(line)
        except json.JSONDecodeError:
            continue # skip malformed records

        # Step 3: Defensive access
        event_type = event.get("type", "").lower()
        status = event.get("status", "").lower()

        # Step 4: Business logic
        if event_type == "issue_created" and status == "success":
            project = event.get("project", "unknown")
            results[project] += 1
    # Step 5: Serialize output
```

```

    return json.dumps(results, indent=2)

# Example usage
raw_input = [
    '{"type": "Issue_Created", "status": "SUCCESS", "project": "JIRA"}',
    '{"type": "issue_created", "status": "success", "project": "CONFLUENCE"}',
    ' malformed json ',
    '',
]
print(process_events(raw_input))

```

3. Why This Code Impresses at Senior Level

- Single-pass O(n) processing
- Defensive against bad input
- Clear variable naming
- Separation of concerns
- Easy to port to Spark (map → filter → groupBy)

4. How You Explain Scaling This to Spark/Databricks

- raw_lines → Spark DataFrame / Dataset
- json.loads → from_json() with schema
- for-loop → map / select
- defaultdict → groupBy + count
- Python logic becomes distributed transformations

5. Typical Atlassian Follow-up Questions

- What happens with billions of rows?
- How do you handle schema evolution?
- How do you monitor bad records?
- How would you test this logic?
- How do you make this idempotent?

6. Final Senior-Level Interview Tips

- State assumptions before coding
- Mention trade-offs explicitly
- Talk about observability (metrics, logs)
- Tie logic back to business outcomes