



Shwetank Singh
GritSetGrow - GSGLearn.com

DATA AND AI

BATCH PROCESSING WITH SPARK

DATAFRAME VS SQL, JOB STRUCTURE

www.gsglearn.com



Batch Processing with Spark --- DataFrame vs SQL, and Job Structure

Goal: Give you a production-ready blueprint for batch jobs in Spark, comparing **DataFrame API vs Spark SQL**, and laying out **job structure**, **idempotency**, **testing**, and **operational patterns**.

Covers: API parity, when to choose which, schema & I/O, joins/aggregations, UDFs, file layout, incremental processing, MERGE/upserts, backfills, orchestration, config knobs, examples in **PySpark** (with Scala/SQL notes).

Table of Contents

1. Mental Model: How Spark Executes
 2. DataFrame vs Spark SQL: Parity, Pros/Cons, and Interop
 3. I/O & Schemas: Read/Write Patterns
 4. Core Transformations & Joins
 5. UDFs vs Built-ins (and Pandas UDFs)
 6. Job Structure Templates
 7. Incremental Processing & Upserts
 8. Backfills & Reprocessing
 9. File Layout on Write
 10. Idempotency, Quality Checks & Error Handling
 11. Configs & Performance Knobs (Batch)
 12. Orchestration: Airflow/Dagster Skeleton
 13. Observability & Logging
 14. Project Skeleton & Testing
 15. End-to-End Examples (DataFrame & SQL)
 16. Cheat Sheet
-

1) Mental Model: How Spark Executes

- Logical Plan → Optimized Plan (Catalyst) → Physical Plan (Tungsten) → Tasks.
- Both **DataFrame** and **Spark SQL** compile to the same Catalyst plan.
- Shuffles happen on **wide** operations (joins, groupBy, window). Keep an eye on them.

```
-- Inspect a query plan (SQL)
EXPLAIN FORMATTED
SELECT c.customer_id, SUM(o.amount) amt
FROM orders o JOIN customers c USING(customer_id)
GROUP BY c.customer_id;
```

2) DataFrame vs Spark SQL: Parity, Pros/Cons, and Interop

Parity: Almost all relational ops exist in both APIs. Interchange via temp views.

```
# Interop
orders_df.createOrReplaceTempView("orders")
spark.sql("SELECT * FROM orders WHERE dt >= '2025-08-01'")
```

When to favor DataFrame API

- Complex programmatic logic, branching, re-usable functions.
- Strong typing (Scala) & IDE refactors.
- Easier unit testing of small transforms.

When to favor SQL

- Pure relational transformations, analysts contributing queries.
- Window-heavy logic is often more readable.
- Leverage optimizer hints (BROADCAST, REPARTITION, etc.).

Team pattern: Keep **business logic** in SQL (views/models), **plumbing** (I/O, params, retries) in DataFrame/PySpark. Mix as needed.

3) I/O & Schemas: Read/Write Patterns

Reads

```
# Parquet/Delta/CSV examples
sales = (spark.read
    .format("parquet")
    .option("mergeSchema", "false") # prefer explicit schemas
    .load("abfss://lake/silver/sales"))

# Explicit schema (recommended)
from pyspark.sql.types import *
schema = StructType([
    StructField("order_id", LongType(), False),
    StructField("customer_id", LongType(), False),
    StructField("dt", DateType(), False),
    StructField("amount", DecimalType(12,2), True),
])
bronze = spark.read.schema(schema).json("abfss://lake/bronze/orders/*.json")
```

Writes

```
# Partitioned write with file sizing
(target_df
    .repartition(200, "dt") # writer parallelism & partitioning
    .sortWithinPartitions("dt", "customer_id") # better encodings
    .write
    .mode("append")
    .option("maxRecordsPerFile", 5_000_000)
    .parquet("abfss://lake/silver/orders"))
```

SQL equivalents

```
CREATE TABLE silver.orders USING PARQUET LOCATION 'abfss://lake/silver/orders';
INSERT INTO silver.orders SELECT * FROM bronze.orders_view;
```

Schema evolution: Prefer **additive** (new nullable columns). Avoid implicit inference in prod.

4) Core Transformations & Joins

```
from pyspark.sql import functions as F
# Filters, projections
clean = (bronze
    .filter(F.col("dt") >= F.lit("2025-08-01"))
    .select("order_id", "customer_id", "dt",
        F.col("amount").cast("decimal(12,2)").alias("amount")))

# Joins
joined = (clean.alias("o")
    .join(customers.alias("c"), "customer_id", "inner")
    .join(products.alias("p"), "sku_id", "left"))

# Aggregations
agg = (joined
    .groupBy("customer_id", F.to_date("dt").alias("d"))
    .agg(F.sum("amount").alias("daily_amount")))
```

SQL

```
WITH clean AS (
    SELECT order_id, customer_id, CAST(amount AS DECIMAL(12,2)) AS amount, dt
    FROM bronze.orders WHERE dt >= DATE '2025-08-01'
)
SELECT c.customer_id, DATE(o.dt) AS d, SUM(o.amount) AS daily_amount
FROM clean o JOIN dim_customers c USING (customer_id)
GROUP BY c.customer_id, DATE(o.dt);
```

5) UDFs vs Built-ins (and Pandas UDFs)

- Prefer **Spark SQL functions** (`functions.*`) → vectorized, optimized, pushdown-friendly.
- **Scala/PySpark UDFs**: easy but slower; break predicate pushdown & codegen.
- **Pandas UDFs** (vectorized) for heavy numeric ops; still avoid if built-ins exist.

```
# Anti-pattern: plain Python UDF for simple math
@F.udf('double')
def bad(x):
    return x * 1.18
```

```
# Better: built-in expr
better = df.select((F.col('amount')*F.lit(1.18)).alias('gross'))
```

6) Job Structure Templates

Template A --- Simple Batch (Bronze → Silver)

```
from argparse import ArgumentParser
from pyspark.sql import functions as F

parser = ArgumentParser()
parser.add_argument('--run-date', required=True) # e.g., 2025-08-22
args = parser.parse_args()

run_date = args.run_date

# 1) Read partition(s)
bronze = spark.read.schema(schema).json(f"abfss://lake/bronze/orders/dt={run_date}/*.json")

# 2) Transform
silver = (bronze
    .filter(F.col('is_valid') == F.lit(True))

    .select('order_id', 'customer_id', 'sku_id', 'amount', F.col('dt').cast('date').alias('dt')))

# 3) Write partitioned
(silver
    .repartition(200, 'dt')
    .write.mode('append')
    .partitionBy('dt')
    .parquet('abfss://lake/silver/orders'))
```

Template B --- SQL-First Batch

```
spark.sql("CREATE TEMP VIEW bronze_orders AS SELECT * FROM
parquet.`abfss://lake/bronze/orders`")
spark.sql("""
INSERT INTO parquet.`abfss://lake/silver/orders`
SELECT order_id, customer_id, sku_id, CAST(dt AS DATE) dt, amount
FROM bronze_orders WHERE is_valid = true
""")
```

Template C --- Multi-stage DAG inside one job

- Stage 1: Load **dimensions** (cached).
- Stage 2: Transform facts with joins.
- Stage 3: Aggregate to daily/monthly tables.
- Stage 4: Optional **optimize/compact** hot partitions.

7) Incremental Processing & Upserts

Incremental by watermark (max dt)

```
from pyspark.sql.functions import col, lit, max as smax
last_dt = (spark.read.parquet('abfss://lake/silver/orders')
            .agg(smax('dt').alias('m')).collect()[0]['m'])
new_data = spark.read.json("abfss://lake/bronze/orders").where(col('dt') >
lit(last_dt))
```

MERGE/Upsert (Delta/Iceberg/Hudi)

```
-- Delta Lake example
MERGE INTO silver.orders AS t
USING tmp_upserts AS s
ON t.order_id = s.order_id
WHEN MATCHED THEN UPDATE SET amount = s.amount, dt = s.dt
WHEN NOT MATCHED THEN INSERT *;
```

Idempotent staging: write tmp_upserts per run_date → MERGE → drop temp.

8) Backfills & Reprocessing

- Drive by **date ranges**; isolate to **target partitions**.
- Use **overwrite by partition** (dynamic) to avoid full table rewrite.

```
(silver
    .write.mode('overwrite')
    .option('partitionOverwriteMode', 'dynamic')
    .parquet('abfss://lake/silver/orders'))
```

- Validate with **row counts** and **checksums** against source.
-

9) File Layout on Write

- Partition primarily by **date/time**; maybe add one moderate-cardinality key.
 - Target **256--1024 MB** files; set **maxRecordsPerFile** and writer parallelism.
 - Periodic **compaction** (Delta `OPTIMIZE`, Iceberg `rewrite_data_files`, Hudi clustering/compaction).
-

10) Idempotency, Quality Checks & Error Handling

- **Idempotency**: deterministic outputs for the same inputs (MERGE + partition overwrite; avoid non-deterministic UDFs).
- **DQ**: Great Expectations/Deequ or SQL asserts.

```
-- Example assert: no null customer_id in silver
SELECT COUNT(*) FROM silver.orders WHERE customer_id IS NULL;
```

- **Quarantine** bad records to an error table with reason.
 - **Retry** only safe sections; avoid duplicate appends.
-

11) Configs & Performance Knobs (Batch)

```
# Parallelism
spark.conf.set('spark.sql.shuffle.partitions', '400') # tune to cluster size/data
volume
# Adaptive Query Execution (AQE)
spark.conf.set('spark.sql.adaptive.enabled', 'true')
# Broadcast joins (threshold in bytes)
spark.conf.set('spark.sql.autoBroadcastJoinThreshold', 104857600) # 100MB
# Skew mitigation (AQE)
spark.conf.set('spark.sql.adaptive.skewJoin.enabled', 'true')
# Caching
facts.cache(); facts.count() # materialize
```

- Prefer **broadcast joins** for small dimension tables.
 - Repartition by **join keys** before heavy joins to reduce skew.
-

12) Orchestration: Airflow/Dagster Skeleton

Airflow DAG (Python)

```
from airflow import DAG
from airflow.providers.apache.spark.operators.spark_submit import SparkSubmitOperator
from datetime import datetime, timedelta

with DAG('orders_daily', start_date=datetime(2025,8,1), schedule='@daily',
catchup=True) as dag:
    bronze_to_silver = SparkSubmitOperator(
        application='/opt/jobs/bronze_to_silver.py',
        name='bronze_to_silver',
        application_args=['--run-date', '{{ ds }}']
    )

    silver_to_gold = SparkSubmitOperator(
        application='/opt/jobs/silver_to_gold.py',
        name='silver_to_gold',
        application_args=['--run-date', '{{ ds }}']
    )

    bronze_to_silver >> silver_to_gold
```

13) Observability & Logging

- **Metrics:** input/output rows, bytes written, task time, shuffle read/write.
 - **Accumulators** for counters (e.g., bad rows quarantined).
 - **Logs:** log params (run date, source paths, target paths, commit IDs).
 - **Lineage:** capture input paths & output table/version.
-

14) Project Skeleton & Testing

```
repo/
  jobs/
    bronze_to_silver.py
    silver_to_gold.py
  conf/
    base.conf
  tests/
    test_transforms.py
  README.md
```

Unit test example (PySpark local)

```
from transforms import compute_daily_amount

def test_daily_amount(spark):
    df = spark.createDataFrame([
        (1, 10, '2025-08-01', 100.0),
        (2, 10, '2025-08-01', 50.0),
    ], 'order_id long, customer_id long, dt string, amount double')
    out = compute_daily_amount(df)
    rows = { (r.customer_id, r.d, r.daily_amount) for r in out.collect() }
    assert rows == { (10, '2025-08-01', 150.0) }
```

15) End-to-End Examples (DataFrame & SQL)

DataFrame pipeline

```
from pyspark.sql import functions as F
bronze = spark.read.json('abfss://lake/bronze/orders/dt=2025-08-22/*.json')
dim = spark.read.parquet('abfss://lake/silver/dim_customers')

silver = (bronze
    .filter('is_valid = true')
    .join(dim, 'customer_id')
    .groupBy('customer_id', F.to_date('dt').alias('d'))
    .agg(F.sum('amount').alias('daily_amount')))

(silver.write
    .mode('append')
```

```
.partitionBy('d')
.parquet('abfss://lake/silver/orders_daily'))
```

SQL pipeline

```
CREATE TEMP VIEW bronze AS SELECT * FROM json.`abfss://lake/bronze/orders/dt=2025-08-22/*.json`;
CREATE TEMP VIEW dim AS SELECT * FROM parquet.`abfss://lake/silver/dim_customers`;

INSERT INTO parquet.`abfss://lake/silver/orders_daily`
SELECT o.customer_id, DATE(o.dt) AS d, SUM(o.amount) AS daily_amount
FROM bronze o JOIN dim USING (customer_id)
WHERE o.is_valid = true
GROUP BY o.customer_id, DATE(o.dt);
```

16) Cheat Sheet

- Prefer **explicit schemas**; avoid inference in production.
- Keep logic **pure** (deterministic) for idempotency; use **MERGE** for upserts.
- Partition by **date**; size files to **256--1024 MB**.
- Turn on **AQE**, tune **shuffle.partitions**, and use **broadcast** for small dims.
- Use **SQL** for relational clarity; **DataFrames** for orchestration/parametrization.
- Validate outputs with **row counts/checksums**; quarantine bad records.

Thank you



Shwetank Singh
GritSetGrow - GSGLearn.com