

ANEETHA CHANDHIRASEKAR

Lead Data Engineer | Data & AI Engineering

Melbourne, VIC | 0415 581 622 | aneetha.chandru108@gmail.com | linkedin.com/in/aneethachandhirasekar/ | Australian Citizen

PROFESSIONAL SUMMARY

Lead Data Engineer with 19+ years of experience designing and delivering enterprise-scale modern data solutions across cloud and on-premises environments. Strong hands-on background in data engineering, lakehouse and data warehouse architecture, ETL/ELT, reusable engineering frameworks, data products, cloud modernisation and large-scale data processing, with deep expertise across Azure Databricks, the Azure data ecosystem. Combines hands-on engineering with solution architecture and technical leadership, leading mid-sized engineering teams and owning multiple complex data engineering initiatives end-to-end from requirements and architecture through build, validation, production deployment and operationalization.

TECHNICAL SKILLS

Languages: Python, PySpark, Advanced SQL, PL/SQL, R

Data Platforms: Azure Databricks, Apache Spark, Delta Lake, Unity Catalog, ADLS Gen2, Azure Data Factory, Snowflake, Oracle

Data Engineering: ETL/ELT, Medallion Architecture, dimensional modelling, Kimball methodology, SCD Type, CDC, batch and incremental processing, data quality, reconciliation

Orchestration & DevOps: Databricks Jobs API, Lakeflow Jobs, Control-M, Azure DevOps, CI/CD, Databricks Asset Bundles

Architecture & Operations: Lakehouse Architecture, modern data warehousing, data products, HLD/LLD, performance and cost optimisation, production engineering, monitoring and observability

BI: Power BI, Tableau, OBIEE, Business Objects

PROFESSIONAL EXPERIENCE

InfoCentric Pty Ltd | Client: Coles Group, Melbourne | Sep 2018 – Present

Lead Data Engineer / Principal Consultant

- Remain hands-on across enterprise data initiatives, building ETL/ELT pipelines, reusable frameworks and production data products with Azure Databricks, PySpark, Python, SQL, Delta Lake and Azure Data Factory.
- Lead system and solution design from requirements through HLD/LLD, choosing data, integration and processing patterns based on scale, performance, reliability, governance and cost.
- Own delivery end-to-end: discovery and stakeholder workshops, solution design, hands-on build, validation, CI/CD, cutover, production deployment and operational handover.
- Lead a mid-sized engineering team through estimation, PI/sprint planning, design and code reviews, engineering standards, production readiness and mentoring while remaining hands-on.
- Partner with Product Owners, BAs, Data Scientists, Architects, Platform/MLOps and business stakeholders; report milestones, risks, dependencies and mitigation to senior management.
- Own production health and optimisation across SLAs, observability, RCA, performance and cloud cost, using compute/query tagging and operational reporting to understand workload usage.
- Contribute to architecture discussions, technical solutioning and RFP activities, defining practical implementation approaches and delivery estimates for enterprise data programs.

KEY INITIATIVES & IMPACT

Enterprise Feature Store Modernisation & Data Products - Coles

- Led enterprise data-product and lakehouse modernisation initiative delivered end-to-end.
- Migrated ~300 enterprise tables from Oracle to Azure Databricks as part of the Feature Store modernisation, using ADLS Gen2, Delta Lake and Medallion Architecture, with ADF supporting enterprise ingestion and data movement where applicable.
- Built batch and incremental ETL/ELT processing in PySpark and Spark SQL, handling complex business transformations across multiple granularities and historical/future data horizons.
- Built a reusable metadata-driven Databricks framework using YAML configuration and status/control tables, integrating Jobs API/Lakeflow Jobs for orchestration and Azure DevOps with Databricks Asset Bundles for controlled deployment.
- Automated migration assurance with a reusable Python regression framework, running source-to-target reconciliation and data-parity checks across 300+ tables.
- Improved production reliability through idempotent processing, watermarks, dependency controls, automated recovery/backfills, SLA monitoring and RCA
- Delivered governed production data products through Unity Catalog and Data Channel, with controlled production deployment, monitoring and operational support.
- Tuned Spark/SQL workloads and Databricks compute for performance and cost, using compute and query tagging to improve workload attribution and resource governance.
- Later led the Feature Store migration to Unity Catalog, covering dependency/impact analysis, validation and production transition while improving governed access, metadata and lineage.
- Improved processing performance by ~30% and reduced cloud cost by ~40% (~\$300K annual savings) by optimising SQL workloads, and right-sizing Databricks SQL/serverless compute.

Range Optimisation - Performance & Architecture Modernisation - Coles

- Investigated the Range Optimisation data flow and Shiny application end-to-end, tracing the main performance and cost bottlenecks back to transformation and solution-design choices.
- Moved data-intensive transformation logic out of R and into Databricks, leaving R focused mainly on modelling and putting scalable data processing in the engineering layer.
- Changed the intermediate design to persist and reuse Delta tables instead of repeatedly recomputing data, improving downstream application efficiency.
- Reworked SQL transformations in Databricks SQL, then moved the workload from Classic to a right-sized Serverless SQL Warehouse based on actual workload and cost characteristics.
- Improved overall processing performance by approximately 30% while reducing compute cost by approximately 15% through architecture correction, transformation optimisation and right-sized serverless compute.

Coles Express Data Product - Azure Data Engineering

- Built and delivered the Coles Express data product end-to-end using Azure Data Factory, ADLS Gen2 and Azure Databricks, from ingestion/data movement through transformation and delivery.
- Built batch, incremental and CDC ETL/ELT pipelines with Azure Data Factory and Databricks, implementing change processing, PySpark/SQL transformations, Medallion Architecture, data-quality checks and source-to-target.
- Took the solution into production with reusable orchestration, Azure DevOps CI/CD, monitoring and operational controls, and supported it through BAU.

Legacy R Processing to Snowflake SQL Modernisation - Coles

- Migrated existing R-based processing and legacy SQL transformation logic to Snowflake SQL using data already available in upstream Snowflake schemas.
- Developed transformation, reconciliation and validation queries to verify business rules and establish data parity between legacy and migrated processing.
- Supported the migration lifecycle from requirements and transformation analysis through regression validation and production transition.

Enterprise Data Warehouse & Dimensional Modelling

- Designed and developed enterprise data warehouse and data-mart solutions using Oracle, ODI/Informatica and PL/SQL, supporting large-scale analytical and reporting workloads.
- Applied Kimball dimensional-modelling principles and SCD Type 1/2 patterns to create historically consistent analytical datasets and reusable warehouse structures.
- Built and optimised ETL transformations, data integration, reconciliation and production processes feeding enterprise BI and operational reporting solutions.

ICON Clinical Research, Chennai | Oct 2013 - Jun 2018

Lead Data Warehouse / BI Developer

- Designed and developed enterprise data warehouses, data marts, ETL/integration solutions and operational reporting using Oracle, ODI, PL/SQL and BI technologies.
- Applied Kimball dimensional modelling and SCD patterns to analytical datasets, supporting reliable historical reporting and enterprise data consumption.
- Delivered enterprise integration and migration initiatives with strong focus on data quality, reconciliation, performance, release management and production reliability.
- Worked with global stakeholders across requirements, development, testing, release and production support while mentoring team members and contributing to engineering standards.

Wipro Technologies, Melbourne & Chennai | Nov 2006 - Aug 2013

Business Intelligence Consultant

- Designed and developed large-scale enterprise data warehouse, ETL and BI solutions using Informatica, Oracle PL/SQL and Business Objects for high-volume transactional and analytical workloads.
- Built ETL transformations and dimensional warehouse structures supporting enterprise data integration, historical processing and downstream business intelligence.
- Contributed to migration and demerger programs, solution design, technical analysis, performance optimisation, risk management and production delivery across distributed teams.
- Coordinated technical delivery across offshore/onshore teams and worked directly with business and technology stakeholders.

CERTIFICATION & EDUCATION

- Databricks Certified Data Engineer Associate
- Bachelor of Engineering - Computer Science Engineering | Anna University, India | 2002-2006 | 79% - First Class with Distinction

RECOGNITION

Above & Beyond Award - InfoCentric (2025) | Customer Satisfaction Award - InfoCentric (2023) | Delivery Excellence Award - ICON Clinical Research | Role Excellence & Team Management Award - Wipro Technologies