

# Modern Data Engineering – End-to-End System Design Guide

Single Source of Truth for Interviews and Real-World Platform Design

## 1. Design Philosophy

- Design for scale, failure, and change from day one.
- Prefer immutable raw data and deterministic transformations.
- Optimize for reprocessing, observability, and cost efficiency.
- Treat data, schema, and quality rules as versioned code.

## 2. Reference Architecture

- Kappa-first architecture: streaming as the default, batch as bounded streaming.
- Medallion pattern: Bronze (raw), Silver (conformed), Gold (serving).
- Separation of concerns: ingestion ≠ transformation ≠ serving.

## 3. Data Ingestion – Source Patterns

File (CSV/JSON): incremental discovery, schema enforcement, quarantine bad records.

RDBMS: CDC preferred; otherwise incremental pulls with partitioned reads.

Kafka/Event Streams: event-time processing, watermarking, deduplication.

## 4. Scalability at Ingestion

- Parallelism: partitions, file splits, micro-batch sizing.
- Backpressure: rate limiting, autoscaling, queue depth monitoring.
- Avoid small files using compaction and correct partition design.

## 5. Schema Management & Data Contracts

- Versioned schemas with backward compatibility rules.
- Explicit nullability, domains, and semantics.
- Machine-readable contracts used in CI/CD and runtime checks.

## 6. Watermarking & Late Data Handling

- Always separate event\_time, ingest\_time, processing\_time.
- Define late arrival tolerance per dataset.
- Watermarks bound state and ensure deterministic aggregates.

## 7. Storage Design

- Columnar formats with ACID support (Delta/Iceberg/Hudi).
- Partition by low-cardinality, query-driven keys.
- Use clustering/Z-ordering for high-cardinality filters.
- Enforce file size standards and compaction.

## 8. Transformations – Mandatory Set

- Type standardization and timezone normalization.
- Null handling with explicit business meaning.
- Deduplication using deterministic keys.
- Reference data enrichment.
- CDC merge logic and SCD Type 1/2 handling.

## 9. CDC & SCD Design

- Capture all change events in Bronze.
- Apply ordered, idempotent merges in Silver.
- Track history where business requires it.
- Support late-arriving changes.

## 10. Scalability in Transformations

- Predicate pushdown and column pruning.
- Broadcast small dimensions.
- Pre-aggregate before joins where possible.
- Skew detection and mitigation strategies.

## 11. Reliability & Reprocessing

- Exactly-once business semantics via idempotent writes.
- Immutable raw data for replay.
- Targeted backfills instead of full recompute.
- Deterministic pipelines.

## 12. Benchmarking Strategy

- Throughput, latency, cost per TB, stability.
- Baseline, stress, and failure injection tests.
- Use benchmarks to drive partitioning and scaling decisions.

## 13. Serving Layer

- Gold tables optimized per consumer.
- BI: pre-aggregations and caching.
- ML: feature freshness and point-in-time correctness.
- APIs: low-latency access paths.

## 14. Monitoring & Observability

- Freshness, completeness, correctness SLOs.
- Lag, volume anomalies, quality drift.
- Cost monitoring and alerts.

## 15. Security & Governance

- RBAC, row/column-level security.
- PII tagging and masking.
- Lineage and auditability.

## 16. Best Practices Cheat Sheet

- Always land raw data immutably.
- Design pipelines to be re-runnable.
- Prefer fewer, larger files.
- Treat data quality as a first-class concern.
- Measure everything.

## 17. Interview Narrative Template

"I start by clarifying SLAs and data correctness requirements. I ingest data immutably into Bronze with contracts and quality checks, apply deterministic CDC/SCD logic in Silver with watermarking for late data, and publish consumer-optimized Gold tables. I ensure scalability via partitioning and parallelism, reliability via idempotent merges and replay, and operational excellence via benchmarking and SLO-driven monitoring."