

# Solution

## 1) Sessionization (30-min gaps)

```
from pyspark.sql import functions as F, Window as W

gap = F.when(F.col("ts") -
F.lag("ts").over(W.partitionBy("user_id").orderBy("ts")) >
F.expr("INTERVAL 30 MINUTES"), 1).otherwise(0)
df1 = (df.orderBy("user_id", "ts")
    .withColumn("new_sess", F.coalesce(gap, F.lit(1)))
    .withColumn("session_id",
F.sum("new_sess").over(W.partitionBy("user_id").orderBy("ts"))
    .rowsBetween(W.unboundedPreceding,
0)))
out = (df1.groupBy("user_id", "session_id")
    .agg(F.min("ts").alias("session_start"),
        F.max("ts").alias("session_end"),
        F.count("*").alias("page_count")))
```

---

## 2) Top-N per group with ties

```
w = W.partitionBy("cust_id").orderBy(F.col("amount").desc())
df_top = (sales
    .withColumn("rnk", F.dense_rank().over(w))
    .filter(F.col("rnk") <= 3))
```

---

## 3) SCD Type-2 (Delta, upserts at scale)

```
from delta.tables import DeltaTable
from pyspark.sql import functions as F

src = source.withColumn("effective_from",
F.col("ingest_ts")).withColumn("effective_to",
F.lit(None).cast("timestamp")).withColumn("is_current", F.lit(True))
```

```

tgt = DeltaTable.forName(spark, "dim_customers")

# Close rows that changed
updates = (src.alias("s").join(tgt.toDF().alias("t"), "customer_id")
    .where(F.sha2(F.to_json("s"), 256) != F.sha2(F.to_json("t"), 256)) # fast change detect on business cols
    .select("t.customer_id"))

tgt.alias("t").merge(
    src.alias("s"), "t.customer_id = s.customer_id AND t.is_current = true"
).whenMatchedUpdate(
    condition="sha2(to_json(s.*),256) <> sha2(to_json(t.*),256)" ,
    set={"effective_to": "s.effective_from", "is_current": "false"}
).whenNotMatchedInsertAll().execute()

```

Tip: Use a **hash of business columns** to avoid expensive row-by-row comparisons.

---

## 4) Incremental de-dup (exactly-once)

**Batch:**

```

# keep only first-seen tx_id across all batches
existing = spark.read.parquet(target_path).select("tx_id").distinct()
new_uniques = incoming.join(existing, "tx_id", "left_anti")
(new_uniques.write.mode("append")).parquet(target_path))

```

**Streaming:**

```

(sdf
    .withWatermark("ts", "2 hours")
    .dropDuplicates(["tx_id"]) # stateful, exactly-once
    semantics per tx_id+watermark
    .writeStream.format("delta").option("checkpointLocation",
    ckpt).start(target_tbl))

```

---

## 5) Skew-safe join

**Broadcast when safe (<= ~10–50MB):**

```
small = F.broadcast(ads.select("ad_id", "attr"))
joined = clicks.join(small, "ad_id", "left")
```

**Salting for hot keys:**

```
from pyspark.sql.functions import rand, floor
salt_n = 20
clicks_s = clicks.withColumn("salt", floor(rand()*salt_n))
ads_s = (ads
    .withColumn("salt", F.explode(F.sequence(F.lit(0),
F.lit(salt_n-1))))))
joined = clicks_s.join(ads_s, ["ad_id", "salt"])
```

Also enable **AQE**: `spark.conf.set("spark.sql.adaptive.enabled", "true")`.

---

## 6) Explode nested JSON (schema-safe)

```
from pyspark.sql import types as T, functions as F

def flatten(df):
    complex_fields = [(f.name, f.dataType) for f in df.schema.fields
                      if isinstance(f.dataType, (T.ArrayType,
T.StructType))]
    while complex_fields:
        name, dtype = complex_fields.pop(0)
        if isinstance(dtype, T.StructType):
            expanded = [F.col(f"{name}.{k}").alias(f"{name}_{k}") for
k in [n.name for n in dtype]]
            df = df.select("*", *expanded).drop(name)
        elif isinstance(dtype, T.ArrayType) and
isinstance(dtype.elementType, T.StructType):
            df = df.withColumn(name, F.explode_outer(name))
            complex_fields = [(f.name, f.dataType) for f in
df.schema.fields]
```

```

        if isinstance(f.dataType, (T.ArrayType,
T.StructType))]
    return df

# enforce provided schema to avoid _corrupt_record
df = spark.read.schema(provided_schema).json(input_path)
flat = flatten(df)

```

---

## 7) Event-time aggregation with late data

```

(sdf
.withWatermark("event_time", "2 hours")
.groupBy(F.window("event_time", "1 hour").alias("win"))
.agg(F.countDistinct("user_id").alias("distinct_users"))
.select("win.start", "win.end", "distinct_users"))

```

---

## 8) Rolling 1-hour avg/stddev with reset (>15-min gap)

```

w_order = W.partitionBy("device_id").orderBy("ts")
gap = (F.col("ts") - F.lag("ts").over(w_order)) > F.expr("INTERVAL 15
MINUTES")
df2 = (df.withColumn("reset", F.when(gap, 1).otherwise(0))
.withColumn("grp",
F.sum("reset").over(w_order.rowsBetween(W.unboundedPreceding, 0))))

```

  

```

w_roll = (W.partitionBy("device_id", "grp")
.orderBy(F.col("ts").cast("long"))
.rangeBetween(-3600, 0)) # 1 hour in seconds (ts cast to
long)

out = df2.select(
"device_id", "ts", "reading",
F.avg("reading").over(w_roll).alias("avg_1h"),
F.stddev_samp("reading").over(w_roll).alias("std_1h")
)

```

---

## 9) File compaction (target ~128MB)

```
spark.conf.set("spark.sql.files.maxRecordsPerFile", 0) # let size
drive splits
# Estimate target partitions by total size / 128MB
input_df = spark.read.parquet(src_path)
target_parts = max(1, int(input_df.rdd.getNumPartitions() / 4)) # quick heuristic
(input_df
    .repartition("dt")           # keep partitioning
    .sortWithinPartitions("dt")   # better packing
    .write.mode("overwrite").partitionBy("dt").parquet(dst_path))
```

On Delta: prefer `OPTIMIZE table ZORDER BY (...)` for ongoing compaction.

---

## 10) CDC merge with deletes (Delta)

```
from delta.tables import DeltaTable
cdc = cdc_df.select("id", "op", "op_ts", "cols...") # op in
('I', 'U', 'D')

tgt = DeltaTable.forName(spark, "sales")
tgt.alias("t").merge(
    cdc.alias("c"), "t.id = c.id"
).whenMatchedUpdate(
    condition="c.op = 'U' AND c.op_ts >= t.op_ts",
    set={"col1": "c.col1", "op_ts": "c.op_ts"}
).whenMatchedDelete(condition="c.op = 'D' AND c.op_ts >= t.op_ts")
.whenNotMatchedInsert(
    condition="c.op = 'I'",
    values={"id": "c.id", "col1": "c.col1", "op_ts": "c.op_ts"}
).execute()
```

Deduplicate CDC first on `(id, op_ts)` with `row_number` to keep the latest per id.

---

## 11) Consecutive daily streaks

```
from pyspark.sql import functions as F, Window as W
d = (logins.select("user_id",
F.to_date("login_date").alias("d")).distinct())
w = W.partitionBy("user_id").orderBy("d")
grp = (F.datediff("d", F.lag("d").over(w)).alias("diff"))
g = (d.withColumn("diff", grp)
    .withColumn("grp",
        F.sum(F.when((F.col("diff") != 1) | F.col("diff").isNull(),
1).otherwise(0))
            .over(w.rowsBetween(W.unboundedPreceding, 0))))
streaks = (g.groupBy("user_id", "grp")
    .agg(F.min("d").alias("start"), F.max("d").alias("end"),
        F.count("*").alias("len")))
result = (streaks.withColumn("r",
    F.row_number().over(W.partitionBy("user_id").orderBy(F.col("len").desc(),
(), F.col("start")))))
    .filter("r=1").drop("r"))
```